

Security Design Doc

Project: Semester 7 - Individual (E2EE Messaging Platform)

Version: 0.1 (2025-11-02)

1. Purpose & Scope

This document consolidates all security-related decisions taken to date and describes the current E2EE protocol design for one-to-one and group messaging, including file sharing. It also lists assumptions, threat model, trust boundaries, and open issues for future iterations.

In scope: Backend services (Gateway, Auth, Messaging), client-server contracts for E2EE, CI/CD security controls.

Out of scope: Full cryptographic proofs, mobile/desktop client UI/UX.

2. High-Level Security Goals

- **Confidentiality:** Server never has plaintext of messages/files; only ciphertext + minimal metadata.
 - **Integrity & Authenticity:** Recipients detect tampering via AEAD auth tags; messages are bound to sender identity.
 - **Forward Secrecy (FS):** Compromise of long-term keys should not reveal past messages (see current status below).
 - **Post-Compromise Safety (PCS):** Limit blast radius of a compromised device via key rotation/re-establishment.
 - **Availability:** Rate-limiting and basic DoS hardening at Gateway; at-least-once delivery with client de-dupe.
 - **Privacy by Design:** Data minimization, purpose limitation, retention, and subject-rights flows.
-

3. Assets, Adversaries, Assumptions

3.1 Assets

- **Message content** (1-to-1 & group).
- **Attachment content** (planned).
- **Long-term identity keys** (client-side).
- **Session/ephemeral keys** (client-side).
- **Account credentials** (email/username + password).
- **Server secrets** (JWT signing key, DB creds).

3.2 Adversaries

- **Network attacker**: eavesdrop/modify traffic; attempts downgrade/miTM.
- **Honest-but-curious server**: can read DB/object store but should learn no plaintext.
- **External attacker**: attempts credential stuffing, token theft, replay.
- **Malicious participant**: valid group member sharing content/keys. (Non-goals: preventing authorized recipients from re-sharing plaintext outside the system.)

3.3 Assumptions

- Clients implement crypto correctly and protect private keys at rest.
- TLS termination is secure; certificates are valid and pinned at the client.
- Users control their devices; device compromise is detectable/recoverable (session revoke).

4. Trust Boundaries & Data Flow

- **Client ↔ Gateway**: Will be TLS-protected HTTP/WS, once deployed on the cloud.
- **Gateway ↔ Services (Auth/Messaging/File)**: internal network (Docker), also TLS-terminable if externalized.
- **Services ↔ Postgres**: private network; credentials via env; least privilege.

5. Implemented Security Controls (Server-Side)

- **Passwords:** Argon2id hash(target ~64 MB memory, p=2, t≈1).
- **Authentication:** Short-lived JWT access (≈15 min) + refresh tokens (≈30 days) bound to a server-side session (revocation on logout/compromise).
- **Sessions:** Store IP as Postgres inet (host only); rotate refresh on use; single-device binding per token.
- **Input Validation:** Strict JSON validation; server rejects plaintext fields where ciphertext is required.
- **Rate Limiting:** Auth ≤ 5 req/sec/IP; Messaging ≤ 60 req/min/user (configurable).
- **Secrets Management:** Environment variables via CI/CD; no secrets in VCS; GH Actions Secrets for sensitive values; Variables for non-sensitive.
- **Logging:** Structured logs; never log secrets or plaintext content; correlation IDs planned.
- **Health & Rollback:** Container healthchecks; immutable image tags :sha-<commit> for rollback.

6. E2EE Protocol - Current Design (MVP)

```
sequenceDiagram
    autonumber
    participant A as Sender (Client A)
    participant G as Gateway
    participant M as Messaging Service
    participant B as Recipient (Client B)
    A->>G: POST /messages/send { envelope(ciphertext, nonce, AAD) }
    G->>M: forward
    M->>B: push over WebSocket
    B->>B: Derive k_enc (X25519 + HKDF) → AEAD decrypt
```

```
sequenceDiagram
    autonumber
    participant UI as Web UI (React)
    participant WASM as WASM Bridge
    participant MC as Msg Client
    participant CC as Crypto Core
    participant KS as Keys Service
    participant MS as Messages Service
    participant DB as Messages DB
```

Note over UI: User presses Send

UI→>WASM: prepareSend with state and plaintext
WASM→>MC: PrepareSend

alt No existing session for this conversation

 MC→>KS: GET /keys/bundle for recipient device
 KS→>MC: PreKeyBundle

 Note over KS,MC: PreKeyBundle fields
 - recipient identity_pub
 - recipient signed_prekey_pub
 - signed_prekey_signature
 - recipient one_time_prekey_pub optional
 - key identifiers for spk and opk

 MC→>CC: InitSession using bundle
 CC→>MC: SessionState and HandshakeMessage

 Note over MC,MS: First message will include Handshake in the header
 Handshake fields sent:
 - sender ephemeral_pub
 - ids of recipient spk and opk used
 - authentication tag or transcript hash

 end

 MC→>CC: Encrypt with SessionState
 CC→>MC: ciphertext and MessageHeader

 Note over MC,MS: Every message header contains ratchet data
 Ratchet h

header fields:
- sender ratchet_dh_pub
- pn previous chain count
- n current chain index
- nonce for AEAD

MC →> WASM: request with header and ciphertext

WASM →> UI: request ready to send

UI →> MS: POST /messages/send with header and ciphertext

MS →> DB: store envelope as ciphertext only

MS →> UI: 200 OK

rect rgb(1,1,1)

Note over MS,DB: Delivery pipeline keeps ciphertext intact

MS →> UI: recipient WS push envelope unchanged

end

UI →> WASM: handleEnvelope with header and ciphertext

WASM →> MC: HandleEnvelope

opt First envelope for this conversation

MC →> CC: AcceptSession using Handshake from header

CC →> MC: SessionState ready for decrypt

end

MC →> CC: Decrypt with header and ciphertext

CC →> MC: plaintext

MC →> WASM: plaintext to UI

6.1 Key Material (Client-Side)

- **Identity Key Pair:** Long-term keypair (Curve25519 for DH; Ed25519 for signatures — optional for MVP). Public identity key is uploaded to server; **private key never leaves client.**
- **Per-Conversation Session Key (MVP):** Established via **X25519** DH between sender and recipient identity keys (or prekeys, see planned); derives a symmetric key via HKDF.
- **Symmetric Cipher: AES-GCM or ChaCha20-Poly1305 (AEAD).** Ciphertext includes **nonce** and **auth tag**.

6.2 One-to-One Messaging (MVP)

Establish session:

1. Sender fetches recipient **public identity key**.
2. Sender computes DH: $K = X25519(\text{sender_priv}, \text{recipient_pub})$.
3. Derive $k_{\text{enc}} = \text{HKDF}(K, \text{context}=\text{"1to1:enc"}, \dots)$.

Encrypt & send:

1. Generate random **nonce** (unique per message).
2. **AAD** binds metadata (e.g., `conversation_id`, `sender_id`, `message_id`, `created_at`).
3. Encrypt payload with AEAD to produce `ciphertext` + `auth_tag`.
4. Send **envelope** to server: `{message_id, conv_id, sender_id, created_at, ciphertext, nonce, auth_tag, content_type, parent_id?}`.

Receive & decrypt:

1. Recipient derives same k_{enc} via DH.
2. Verify AEAD tag over **ciphertext + AAD**; on success, return plaintext.

Delivery semantics: at-least-once; clients **de-duplicate** via `message_id`.

7. Message Envelope (Server-Visible)

```
{  
  "message_id": "uuid-v4",  
  "conversation_id": "uuid-v4", // or group_id  
  "sender_id": "uuid-v4",  
  "created_at": "RFC3339",  
  "ciphertext": "base64",  
  "nonce": "base64",  
  "auth_tag": "base64",  
  "content_type": "text|file-pointer|...",  
  "parent_id": "uuid-v4|null" // optional reply threading  
}
```

8. Secure Development Lifecycle (SDLC) Controls

- **Linters/Static Analysis:** `golangci-lint` (gofumpt, gocyclo, revive, unparam, misspell).
 - **CI:** `go vet`, `go test`, build per service; fail-fast.
 - **Supply Chain:** Images built in CI; pushed to GHCR; deploys pin `:sha-<commit>`.
 - **Environments:** `.env` templating via Actions; secrets injected only at deploy; no secrets in logs.
 - **Observability (MVP):** Health endpoints + container healthchecks; planned centralized logs/metrics.
-
-